

AI-Based Early Detection of Alzheimer's Disease Using Medical images

A. Assignment Information

Particular	Details
Department:	Computer Science Engineering
Programme:	B .Tech/B.E
Course Code & Course Name	CSA1708 & Artificial Intelligence
Academic Year / Batch	2026-2027/2 nd year
Faculty Name	Dr .R .Jeena
Assignment Title	AI-Based Early Detection of Alzheimer's Disease Using Medical images
Date of Issue	1.09.2026
Date of Submission	2.09.2026
Maximum Marks	100
Course Outcome(s) – CO	C04:Apply machine learning and deep learning techniques to solve real-world biomedical and healthcare data problems.
Bloom's Taxonomy Level	LA-Analyze,L5-evaluate,L6-Create
SDG Mapping (SDG 1-17)	SDG 3:Good Health and Well-Being
Industry / Societal Relevance	Early and low-cost detection of Alzheimer's Disease using AI-assisted analysis of MRI/PET scans support radiologists, reduces diagnostic delay, and improves patient outcomes-a growing need given the rising global burden of dementia.

B. Assignment Problem / Challenge

Faculty shall provide a **realistic engineering/scientific/computing problem, case, challenge or design requirement** related to the Course Outcome.

The assignment should preferably require students to:

- Apply course concepts rather than reproduce theory.
- Identify and analyze the problem.
- Consider requirements and constraints.
- Develop or compare alternative solutions where appropriate.

- Interpret results.
- Make and justify engineering decisions.
- Consider appropriate technical, economic, environmental, societal, ethical, safety or sustainability factors wherever relevant.

C. Problem Statement

Problem / Case / Design Challenge:

[Faculty to provide the detailed assignment problem here.]

D. Requirements and Constraints

Students shall identify or be provided with relevant requirements and constraints. Examples

include:

- Functional requirements
- Performance requirements
- Technical constraints
- Cost, Time
- Resources
- Safety
- Reliability
- Sustainability
- Environmental and Ethical considerations
- Standards/codes
- User requirements
- Manufacturability
- Power/energy
- Security/privacy

Only constraints relevant to the particular course/problem need to be considered.

E. Student Work

1. Problem Understanding and Formulation

Explain:

- What is the problem?
- What are the expected outcomes?
- What information/data is available?
- What assumptions are made?

- What constraints must be satisfied?

2. Application of Course Knowledge

Identify and apply the relevant:

- Principles
- Theories
- Mathematical models
- Algorithms
- Engineering concepts
- Scientific concepts
- Design methodologies

Show necessary calculations, derivations or reasoning.

3. Solution / Design / Methodology

Develop the proposed solution. Depending upon the nature of the course, this may include:

- Design
- Algorithm
- Model
- Circuit
- Architecture
- Experiment
- Process
- Prototype
- Simulation
- Case analysis
- Data analysis
- Mathematical solution

Where appropriate, consider **more than one possible solution**.

4. Use of Modern Tools

Use appropriate tools wherever relevant.

Examples:

CSE/IT: Python, MATLAB, TensorFlow, cloud platforms, Git, simulation tools

ECE/EEE: MATLAB/Simulink, Cadence, Synopsys, Vivado, Multisim, LTspice

Mechanical: ANSYS, SolidWorks, CATIA, MATLAB

Civil: STAAD.Pro, ETABS, AutoCAD, Revit

Biomedical: MATLAB, LabVIEW, signal/image-processing tools

Biotechnology: Bioinformatics/data-analysis/simulation tools

Students shall provide appropriate evidence of tool usage and outputs.

5. Results and Validation

Present the results using appropriate:

- Tables
- Graphs
- Simulation outputs
- Experimental observations
- Screenshots
- Performance metrics
- Statistical analysis
- Test results

Validate whether the proposed solution satisfies the stated requirements.

6. Analysis and Engineering Decision

Interpret the results.

Where applicable:

- Compare alternatives.
- Identify advantages and limitations.
- Discuss trade-offs.
- Explain why the final solution was selected.
- Justify the decision using quantitative or qualitative evidence.

7. Broader Considerations

Where relevant to the problem, discuss the impact of the proposed solution with respect to:

- Sustainability
- Environment
- Society
- Safety
- Ethics
- Economics
- Accessibility

- Professional responsibility

8. Conclusion

Summarize:

- Proposed solution
- Major findings
- Achievement of requirements
- Limitations
- Possible improvements

9. Student Reflection

- What did you learn from this assignment beyond what was directly taught in the classroom?
- If you were given additional time/resources, what would you improve and why?

10. References

Use an appropriate citation format and acknowledge all external sources, datasets, standards, software and AI-assisted tools used.

F. Common Assessment Rubric

Assessment Criterion	Marks
Problem understanding & formulation	10
Application of course/domain knowledge	20
Solution methodology / design / implementation	20
Use of appropriate modern tools / techniques	10
Results, testing & validation	15
Analysis, trade-offs & justification	15
Broader considerations / professional responsibility, where applicable	5
Technical documentation & reflection	5
TOTAL	100

G. CO–PO–Assessment Mapping

Assessment Component	CO	PO(s)	Bloom's Level	Marks
Problem formulation	CO5	1,2,3,4,5,8,9,11,12	L3/L4	10
Application of knowledge	CO5	1,2,3,4,5,8,9,11,12	L3/L4	20
Solution / Design / Implementation	CO5	1,2,3,4,5,8,9,11,12	L4/L5/L6	20
Modern tool usage	CO5	1,2,3,4,5,8,9,11,12	L3/L4	10
Validation	CO5	1,2,3,4,5,8,9,11,12	L4/L5	15
Analysis & justification	CO5	1,2,3,4,5,8,9,11,12	L4/L5	15
Broader considerations	CO5	1,2,3,4,5,8,9,11,12	L4/L5	5
Documentation & reflection	CO5	1,2,3,4,5,8,9,11,12	L3/L4	5

Note: Faculty shall map only the POs and Bloom's levels genuinely assessed by the particular assignment. Every assignment is **not required to map to every PO**.

H. Faculty Evaluation & Continuous

Improvement CO Attainment

Performance Level	Criterion
Level 3	$\geq 70\%$
Level 2	60–69%
Level 1	50–59%
Level 0	$< 50\%$

Target CO Attainment: $\geq 70\%$ _____

Actual CO Attainment: **90%**_____

Faculty Analysis

Areas in which students performed well:
Completed academic activities on time.

Needs improvement in accuracy.
Give regular practice and feedback.

Areas requiring improvement:

**Corrective / Improvement Action: Conduct
periodic assessments and reviews.**

Action to be implemented in the next offering:

Documents to be Maintained

The department/course faculty shall maintain:

1. Assignment question/problem statement
2. CO–PO and Bloom's mapping
3. Assessment rubric
4. Student submissions
5. Tool/simulation/experimental evidence, where applicable
6. Evaluated scripts/reports
7. Marks and rubric-wise assessment
8. CO attainment analysis
9. Sample student work representing different performance levels
10. Faculty analysis and continuous-improvement action

A. Executive Summary

The e-commerce platform under study distributes requests across four backend servers of unequal capacity using plain round-robin forwarding, which leaves weaker servers saturated during peak traffic while stronger servers idle – breaching the platform's 2-second SLA precisely when it matters most, during flash sales. This report designs, simulates and evaluates four load-balancing strategies of increasing sophistication – Round Robin (RR), Least Connections (LC), Weighted Round Robin (WRR) and an adaptive Least Response Time (LRT) balancer – using a discrete-event simulation built in Python/SimPy, validated against closed-form queueing-theory predictions and repeated across five random seeds under both normal and flash-sale-style peak traffic.

The results are unambiguous: Round Robin is not merely suboptimal but a reliability risk, with average response time increasing roughly 13x (to ≈ 14.9 s) and 95th-percentile latency increasing roughly 14x (to ≈ 56.4 s) under a 2x traffic surge. Weighted Round Robin and Least Connections both remain stable under peak load, and the adaptive Least Response Time algorithm performs best on every metric in both scenarios – cutting average response time by $\approx 43\%$ and tail latency by $\approx 51\%$ relative to Round Robin under normal load, while requiring no hardware changes and no application-code modification.

B. Division of Work Between Team Members

This assignment was carried out jointly. The table below records who led which part of the work, in line with the documentation requirement in Section E.10 and the professional-responsibility expectation in Section E.7. Both members reviewed and agreed on the final analysis and conclusions before submission.

Task	Lead
Problem understanding, requirements & constraints (Sections C, D)	Member 1
Algorithm design & simulation implementation in Python/SimPy (Section E.2, E.3)	Member 2
Running experiments, statistical analysis, chart generation (Section E.5)	Member 1 & 2 (jointly)
Architecture diagram and system design write-up (Section E.3)	Member 1
Engineering analysis, trade-off discussion and final recommendation (Section E.6)	Member 2
Broader considerations – sustainability, economics, reliability (Section E.7)	Member 1
Risk analysis, deployment plan and monitoring plan (Sections F, G, H)	Member 1 & 2 (jointly)
Report compilation, formatting, references and appendix (Sections E.10, Appendix)	Member 2

C. Problem Statement

An online e-commerce platform experiences highly variable traffic – moderate load through the day, sharp surges during flash sales, and occasional regional spikes around festival periods. The platform runs its backend on four application server instances of unequal capacity: two standard instances and two higher-capacity instances (roughly double and triple the processing power of a standard instance). Under the platform's current setup, incoming HTTP requests are forwarded to backend servers without any intelligent distribution strategy, which has resulted in uneven server utilization, occasional request queuing on smaller instances while larger instances remain under-used, and response times that periodically breach the internal Service Level Agreement (SLA) of 2 seconds average latency.

The task is to design, simulate and evaluate a request-distribution (load-balancing) strategy that reduces average and tail response time, keeps server utilization balanced in proportion to each server's actual capacity, and continues to behave acceptably not only under normal traffic but also under a flash-sale-style traffic surge – without requiring additional hardware.

D. Requirements and Constraints

Functional Requirements

- Every incoming client request must be routed to exactly one of the four backend servers.
- The routing decision must be made in real time, without introducing noticeable overhead of its own.
- The solution must account for the fact that the four servers have unequal processing capacity.
- The solution must remain functionally correct under both normal traffic and a 2x flash-sale traffic surge.

Performance Requirements

- Average response time across all requests should not exceed the 2-second internal SLA under normal load.
- 95th-percentile (tail) response time should be minimised, since tail latency is what most affects perceived checkout experience.
- Throughput (requests served per second) must not degrade compared to the current unbalanced setup.

- Performance should degrade gracefully, not catastrophically, when traffic doubles during a flash sale.

Technical Constraints

- No new hardware or additional server instances may be purchased (cost constraint).
- The solution must be implementable at the load-balancer layer only, without modifying application code.
- The chosen algorithm must be simple enough to reason about and to tune operationally (maintainability constraint).

Other Considerations

- Reliability – a server that becomes slow or unavailable should not continue to receive a full share of traffic.
- Sustainability – better utilisation of existing servers reduces the need for over-provisioning, lowering energy consumption.
- Security / privacy – session and cart data handling is unaffected by the choice of balancing algorithm, but must not be broken by it.

E. Student Work

1. Problem Understanding and Formulation

What is the problem? Requests are currently distributed to four heterogeneous backend servers without regard to their differing capacities, causing uneven load, queuing delays on weaker servers, and SLA breaches during peak traffic.

Expected outcomes: A load-balancing strategy (and a working simulation of it) that lowers average and 95th-percentile response time, keeps utilisation proportional to server capacity, and respects the SLA of 2 seconds under normal load, evaluated under both normal and flash-sale traffic, without adding servers.

Data / information available:

- Relative capacity of each server, expressed as integer weights (1, 1, 2, 3).
- An estimate of normal request arrival pattern (roughly one request every 0.6 seconds) and a flash-sale pattern (roughly one request every 0.3 seconds).
- An estimate of average per-request processing time on a standard-capacity server (1.2 seconds).

Assumptions made:

- Request inter-arrival times and per-request service times follow an exponential distribution, a standard and reasonable approximation for independent web traffic.

- A server's effective service rate scales linearly with its assigned capacity weight.
- Requests are stateless at the routing layer (session persistence is handled separately and is outside the scope of this comparison).
- The flash-sale scenario is approximated as a sustained doubling of arrival rate rather than a short, extreme spike.

Constraints to satisfy: No new hardware, SLA of 2 seconds average response time under normal load, graceful degradation under peak load, and a routing decision computable in $O(1)$ – $O(n)$ time per request for $n = 4$ servers.

2. Application of Course Knowledge

This problem is a classic application of queuing theory and distributed-systems load-distribution algorithms covered in the course. Four candidate algorithms were selected for comparison, extending the original three-way comparison with a fourth, adaptive algorithm:

- **Round Robin (RR)** – requests are sent to servers in strict cyclic order, regardless of load or capacity.
- **Least Connections (LC)** – each request is sent to whichever server currently has the fewest active/queued requests.
- **Weighted Round Robin (WRR)** – requests are distributed cyclically but in proportion to each server's declared capacity weight, using the smooth WRR pointer algorithm (as used internally by NGINX).
- **Least Response Time (LRT)** – each request is sent to the server with the lowest exponentially-weighted moving average (EWMA) of recent response times, adjusted for its current queue length; this lets the balancer adapt to real, observed performance rather than relying only on a declared weight.

Time and space complexity of the four algorithms

Algorithm	Time / decision	Space	Adapts to live load?	Needs known weights?
Round Robin	$O(1)$	$O(1)$	No	No
Least Connections	$O(n)$	$O(n)$	Yes	No
Weighted Round Robin	$O(n)$	$O(n)$	No	Yes
Least Response Time	$O(n)$	$O(n)$	Yes	No (learns online)

Alternative Algorithms Considered and Rejected

Before settling on the four algorithms above, three additional well-known strategies were considered and deliberately excluded, for reasons documented here rather than left implicit:

- **Random Assignment** – each request routed to a uniformly random server. It shares Round Robin's blindness to capacity and load, and a capacity-weighted variant would collapse into Weighted Round Robin's continuous analogue, which is already covered; it was therefore judged to add no distinct data point.
- **Source-IP Hash / Consistent Hashing** – routes a request based on a hash of the client's IP address, mainly to give session affinity (“sticky sessions”) or cache locality. Assumption (c) in Section E.1 states that requests are stateless at the routing layer for this comparison, so the one property this family of algorithms optimises for is out of scope here; it remains relevant if session affinity becomes a requirement later.
- **Power-of-Two-Choices (P2C)** – sample two servers at random and route to whichever of the two currently has fewer in-flight requests. P2C is a genuinely strong middle ground – it approaches Least Connections' quality with $O(1)$ rather than $O(n)$ bookkeeping – but with only $n = 4$ servers the $O(n)$ cost of full Least Connections is already negligible, so P2C's main advantage does not materialise at this scale. It is flagged in Section E.7 as a stronger candidate if the server pool grows substantially.

Queueing-Theory Foundations

Two standard results from queueing theory were used to reason about expected behaviour before simulating:

- **Little's Law:** $L = \lambda \times W$, where L is the average number of requests in the system, λ is the arrival rate, and W is the average time a request spends in the system.
- **Utilisation:** $\rho = \lambda / \mu$, where μ is the service rate of a server (or pool of servers).
- **M/M/1 mean waiting time:** for a single server modelled as an M/M/1 queue with arrival rate λ and service rate μ , the expected time a request spends in the system is $W = 1 / (\mu - \lambda)$, valid only while $\rho = \lambda / \mu < 1$.

Aggregate pool capacity is $\mu_{\text{total}} = \Sigma(\text{weight}_i / 1.2) = 7 / 1.2 \approx 5.833$ requests/second (weights sum to $1+1+2+3 = 7$). Under normal load ($\lambda \approx 1.667$ req/s), aggregate utilisation is $\rho \approx 0.286$ (about 29%), so the pool clearly has enough raw capacity – confirming the SLA breaches in Section C are a distribution problem, not a capacity problem.

Closed-Form Sanity Check: RR vs WRR Under Normal Load

Before trusting the simulation, its two extreme algorithms were checked against the M/M/1 formula above, treating each server as an independent single-server queue.

Round Robin sends each server an equal $1/4$ share regardless of weight, so $\lambda_i = 1.667/4 \approx 0.417$ req/s for every server. With $\mu_1 = \mu_2 \approx 0.833$, $\mu_3 \approx 1.667$ and $\mu_4 \approx 2.5$ req/s, the M/M/1 formula gives $W_1 = W_2 \approx 2.40$ s, $W_3 \approx 0.80$ s and $W_4 \approx 0.48$ s. Averaging across all four servers (each carrying an equal share of requests) gives a predicted overall average of ≈ 1.52 s – the same order of magnitude as, though somewhat higher than, the 1.144 s the simulation measured. The gap is expected: the closed-form M/M/1 result is an asymptotic steady-state average, while the simulation runs for a large but finite 2,000 seconds and captures transient correlation effects that a hand calculation ignores; the agreement in order of magnitude and in relative ranking is exactly what a sanity check is meant to confirm.

Weighted Round Robin sends each server $\lambda_i = \lambda_{\text{total}} \times (w_i/7)$, which by construction makes $\rho_i = \lambda_i/\mu_i$ identical for every server regardless of weight ($\rho \approx 0.286$ for all four, matching the pool-wide utilisation computed above). This is the defining theoretical property of capacity-proportional load balancing: WRR does not just move the average down, it equalises risk across the fleet so no single server is closer to saturation than any other. Plugging the four per-server W values into a weighted average (weighted by each server's request share) predicts an overall average of ≈ 0.96 s, again the same order of magnitude as the simulated 0.794 s.

Under Round Robin specifically, each server receives an equal $1/4$ share of traffic regardless of weight. At peak load ($\lambda \approx 3.333$ req/s), each of the two weight-1 servers individually receives ≈ 0.833 req/s, but each can only serve at $\mu_i = 1/1.2 \approx 0.833$ req/s – giving a local utilisation $\rho_{\text{local}} \approx 1.0$. Queueing theory predicts that as $\rho \rightarrow 1$, waiting time grows without bound; this is exactly the runaway response time that Section E.5 measures for RR under peak load, and it is the central theoretical reason RR is unsuitable for this heterogeneous server pool.

3. Solution / Design / Methodology

The proposed solution keeps the existing four servers and inserts a load balancer capable of running any of the four candidate strategies in front of them, as shown in Figure 1.




Figure 1: Load-balanced architecture for the e-commerce backend.

To evaluate the four candidate algorithms before recommending one for production, a discrete-event simulation was built in Python using the SimPy library. The simulation models:

- An arrival process generating requests with exponentially distributed inter-arrival times (mean 0.6 s for normal load, 0.3 s for the flash-sale scenario).
- Four server resources, each able to process one request at a time, with per-request service time drawn from an exponential distribution whose mean is inversely proportional to the server's capacity weight.
- A pluggable routing function implementing RR, LC, WRR or LRT, selected per simulation run.
- Instrumentation that records the response time (queue wait + processing time) of every completed request, and, for LRT, an EWMA tracker per server updated after every completed request.

Selection Logic – Pseudocode

For clarity, the routing rule for each algorithm is restated below as pseudocode (the working Python implementation is reproduced in full in Appendix A):

```

function PICK_RR():
    i ← pointer mod 4
    pointer ← pointer + 1
    return server[i]

function PICK_LC():
    return argmin_i( queue_len[i] )
  
```

```

function PICK_WRR():
    for i in 1..4: current[i] ← current[i] + weight[i]
    best ← argmax_i( current[i] )
  
```

```
current[best] ← current[best] – sum(weight)
return server[best]
```

```
function PICK_LRT():
    return argmin_i( ewma_rt[i] × (1 + queue_len[i]) )
```

```
function ON_COMPLETE(i, response_time):
    ewma_rt[i] ← 0.2 × response_time + 0.8 × ewma_rt[i]
```

Testing Methodology

To validate the simulation itself (not just the algorithms), each routing function was unit-tested independently of the discrete-event engine – for example, feeding the WRR selector a known weight sequence [1,1,2,3] and checking that, over one full cycle of 7 selections, each server is chosen exactly as many times as its weight; and confirming the RR selector visits all four servers exactly once every four calls. Each algorithm was then simulated over 2,000 seconds of simulated time, independently repeated with five different random seeds per scenario, so that reported figures are a mean across runs with a standard deviation showing run-to-run variability rather than a single, potentially unrepresentative run.

4. Use of Modern Tools

- **Python 3** – simulation and analysis logic.
- **SimPy** – discrete-event simulation framework used to model server resources, request arrivals and queuing behaviour realistically.
- **Statistics module** – computing mean, standard deviation and 95th-percentile response times across five repeated runs per scenario.
- **Matplotlib** – generating the comparison charts (Figures 2–4) and the architecture diagram (Figure 1).
- **Git** – used by the team to version-control the simulation code and coordinate changes between both members.

5. Results and Validation

Each of the four algorithms was simulated five times (different random seeds) under two traffic scenarios: normal load and a flash-sale-style peak load at roughly twice the arrival rate. Reported values are the mean across the five runs, with standard deviation in parentheses. Full per-seed figures are given in Appendix B.

Table 1: Normal load (mean inter-arrival 0.6 s)

Algorithm	Avg. Response Time	95th %ile Response Time	Throughput
-----------	--------------------	-------------------------	------------

Round Robin	1.144 (± 0.039) s	4.090 (± 0.192) s	1.668 req/s
Least Connections	1.009 (± 0.019) s	3.293 (± 0.088) s	1.674 req/s
Weighted Round Robin	0.794 (± 0.028) s	2.598 (± 0.076) s	1.676 req/s
Least Response Time	0.652 (± 0.011) s	2.005 (± 0.047) s	1.666 req/s




Figure 2: Response time by algorithm under normal load, mean $\pm$ std across 5 runs.

Table 2: Peak / flash-sale load (mean inter-arrival 0.3 s, $\approx 2\times$ normal arrival rate)

Algorithm	Avg. Response Time	95th %ile Response Time	Throughput
Round Robin	14.925 (± 2.497) s	56.433 (± 8.296) s	3.284 req/s
Least Connections	1.117 (± 0.019) s	3.804 (± 0.098) s	3.317 req/s
Weighted Round Robin	1.128 (± 0.017) s	3.677 (± 0.130) s	3.315 req/s
Least Response Time	0.964 (± 0.035) s	2.898 (± 0.117) s	3.319 req/s




Figure 3: Response time by algorithm under peak load, mean $\pm$ std across 5 runs. Note the y-axis scale – Round Robin's tail latency is over an order of magnitude worse than the other three algorithms.

Validation against the requirements in Section D: under normal load, all four algorithms keep average response time under the 2-second SLA; under peak load, Round Robin breaks the SLA catastrophically (average ≈ 14.9 s, 95th

percentile ≈ 56.4 s) while Least Connections, Weighted Round Robin and Least Response Time all remain within a few seconds, with Least Response Time performing best on every metric in both scenarios.

Traffic-Share Fairness: Does Each Server Get Its Proportional Share?

Response time alone does not show whether a server is being used in proportion to its declared capacity. The simulation was therefore instrumented to also record what fraction of total requests each server actually received, and this was compared against the capacity-proportional ideal ($1/7$, $1/7$, $2/7$, $3/7$).




Figure 4: Actual traffic share per server vs. the capacity-proportional ideal, under normal load.

Server (weight)	Ideal share	RR	LC	WRR	LRT
Server 1 (w=1)	14.3%	25.0%	34.6%	14.3%	4.0%
Server 2 (w=1)	14.3%	25.0%	26.8%	14.3%	3.3%
Server 3 (w=2)	28.6%	25.0%	26.1%	28.6%	30.1%
Server 4 (w=3)	42.9%	25.0%	12.6%	42.9%	62.7%

This reveals a subtlety the response-time tables alone do not: WRR matches the ideal capacity-proportional split almost exactly (by construction), whereas Least Connections is measurably unfair – it over-sends to Server 1 (34.6% vs an ideal 14.3%) simply because a weight-1 server's queue empties out faster in absolute request count, even though each of those requests takes it longer to finish. Least Response Time is unfair in the opposite direction: it deliberately concentrates 62.7% of traffic on the strongest server (weight 3, ideal 42.9%) and starves the weakest two servers, because it is optimising purely for measured latency rather than for a proportional split. This is a genuine engineering trade-off, discussed further in Section E.6: LRT's fairness deviation is precisely why it achieves the best raw latency numbers, but it also means the weight-1 servers sit comparatively idle – a form of under-utilisation that a purely capacity-fairness-driven operator might find undesirable even though the user-facing SLA numbers improve.

6. Analysis and Engineering Decision

Under normal load, the ranking from worst to best is RR, LC, WRR, LRT – each successive algorithm uses more information (load state, known capacity, or measured performance) and each reduces both average and tail latency accordingly. LRT's advantage over WRR under normal load (0.652 s vs 0.794 s average) shows that reacting to real observed performance can outperform even a correctly-configured static weighting.

Under peak load, the picture changes sharply. Round Robin's average response time jumps roughly 13x, from about 1.14 s to about 14.93 s, and its tail latency jumps roughly 14x, from about 4.09 s to about 56.43 s. This matches the theoretical prediction in Section E.2: at peak arrival rate, RR pushes the two weight-1 servers to a local utilisation $\rho \approx 1.0$, and queueing delay explodes as utilisation approaches saturation. The other three algorithms remain stable at peak load because they actively account for either current load (LC, LRT) or known capacity (WRR), so no single server is driven to saturation.

Using Little's Law ($L = \lambda W$) on the normal-load average response times gives an average of roughly 1.91 in-flight requests for RR, 1.68 for LC, 1.32 for WRR and 1.09 for LRT – a concrete, quantitative confirmation that more load-aware algorithms keep fewer requests waiting in the system at any given moment.

Section E.5's fairness analysis adds a necessary caveat to this ranking: LRT's best-in-class latency numbers come partly from deliberately under-using the two weakest servers (a combined 7.3% of traffic against an ideal 28.6%) rather than from a uniformly better scheduling rule. This is an acceptable, even desirable, trade-off for a latency-SLA-driven business like e-commerce checkout, but it does mean the weight-1 servers are not being “used well” in a raw utilisation sense – a point worth flagging explicitly to infrastructure stakeholders rather than leaving implicit in a single headline latency number.

Engineering decision: Least Response Time is recommended for production. It performed best on every metric in both scenarios, and – unlike Weighted Round Robin – it does not depend on manually-declared, potentially stale capacity weights; it learns server performance from live traffic. Its cost is slightly higher implementation and monitoring complexity than WRR (it must track and update a per-server EWMA) and a deliberate skew away from proportional utilisation, both of which the team judged acceptable given the clear and consistent latency gain, especially the far greater robustness under flash-sale conditions. Weighted Round Robin remains a reasonable, simpler fallback if the team prefers predictable, capacity-proportional utilisation over

the last few hundred milliseconds of latency, and Least Connections is an acceptable minimum-viable improvement over the current Round-Robin-equivalent baseline. Round Robin itself should be retired for this heterogeneous server pool – the peak-load results show it is not just suboptimal but a genuine reliability risk during exactly the high-traffic events (flash sales) when the platform can least afford downtime.

7. Broader Considerations

- **Sustainability / Environment:** better-balanced utilisation of the existing four servers reduces the likelihood of provisioning additional server instances purely to compensate for poor distribution, which in turn reduces energy consumption – directly supporting SDG 9 and SDG 12.
- **Economics:** as a rough order-of-magnitude illustration, if avoiding Round Robin's peak-load breakdown removes the need for even one extra always-on server instance that would otherwise be provisioned purely to survive flash sales, and that instance costs on the order of a few hundred dollars a month to run, the saving compounds to a meaningful annual figure for no added hardware cost – the entire improvement is a configuration-level change.
- **Reliability:** in production, LRT should be paired with basic health checks so that a server which becomes slow or fails is not just deprioritised by its EWMA but actively removed from rotation (see Section F).
- **Accessibility / society:** lower tail latency disproportionately helps users on slower or less stable connections, who are most affected by worst-case response times during checkout, and are hit hardest of all under an RR-style collapse during a sale.
- **Scalability:** if the server pool grows well beyond four instances, the $O(n)$ per-request cost of LC and LRT should be re-examined; Power-of-Two-Choices (Section E.2) becomes an attractive upgrade path at that point.
- **Professional responsibility:** the recommendation is based on simulated evidence with clearly stated assumptions, repeated across five random seeds per scenario to show it is not a one-off result, but those assumptions (arrival pattern, service-time distribution, the specific 2x peak multiplier) should still be re-validated against real production traffic logs before final sign-off.

8. Conclusion

Four load-balancing strategies – Round Robin, Least Connections, Weighted Round Robin and Least Response Time – were simulated, each over five random seeds, under both normal and flash-sale-style peak traffic for a four-server, heterogeneous-capacity e-commerce backend, and cross-checked against closed-form queueing-theory

predictions. Under normal load, response time improved progressively from Round Robin through to Least Response Time, which reduced average response time by roughly 43% and 95th-percentile response time by roughly 51% compared to Round Robin. Under peak load, the difference became qualitative rather than incremental: Round Robin's response times increased by more than an order of magnitude and breached the SLA outright, while the other three algorithms remained stable, with Least Response Time again performing best – albeit at the cost of a measurable fairness skew away from proportional server utilisation. All requirements identified in Section D were satisfied by the recommended Least Response Time strategy, achieved through a configuration-level change requiring no additional hardware. The main limitation of this study is that it relies on simulated, exponentially-distributed traffic rather than a production traffic trace, and models the flash-sale scenario as a sustained rate doubling rather than a short, extreme burst; both simplifications should be revisited with real traffic data before production rollout.

9. References

- G. Coulouris, J. Dollimore, T. Kindberg and G. Blair, a standard graduate textbook on distributed systems concepts and design, consulted for background on load distribution and queuing behaviour in distributed systems.
- SimPy official documentation, consulted for discrete-event simulation constructs (Environment, Resource, Process) used to build the simulation.
- NGINX official documentation on HTTP load-balancing methods, consulted for the smooth Weighted Round Robin selection rule and for background on least-response-time-style balancing used in this simulation.
- Standard queueing-theory references for Little's Law and utilisation ($\rho = \lambda/\mu$), including J.D.C. Little's original 1961 derivation, used in Sections E.2 and E.6 to reason about expected system behaviour analytically before simulating.
- AWS Elastic Load Balancing documentation, consulted for how least-outstanding-requests and round-robin routing policies are exposed and configured in a real commercial load balancer.
- Kubernetes documentation on Service and Ingress load balancing, consulted for how capacity-aware routing and health-check-based endpoint removal are implemented in a container-orchestration setting.
- AI-assisted tool: Claude (Anthropic) was used to help structure the simulation code, generate the comparison charts and architecture diagram, and organise this report; all analysis, interpretation and the final engineering decision were reviewed by both team members.

F. Risk Analysis and Mitigation (FMEA)

A lightweight Failure Mode and Effects Analysis (FMEA) was carried out on the recommended Least Response Time design, to surface risks that a pure performance comparison does not capture.

Failure mode	Likelihood	Impact	Mitigation
EWMA update bug silently skews routing toward one server	Low	High	Unit tests on the routing function (Section E.3); alert if any server's traffic share deviates from its rolling average by more than a fixed threshold.
Health check false negative removes a healthy server from rotation	Medium	Medium	Require multiple consecutive failed checks before removal; automatic re-admission with a brief traffic ramp-up rather than an instant full share.
Load balancer restart resets all EWMA state to the base service time	Medium	Medium	Persist EWMA state across restarts where possible; otherwise accept a brief (seconds-scale) re-learning window and monitor it explicitly.
Traffic surge exceeds the assumed 2x flash-sale multiplier	Medium	High	Load-test beyond 2x in staging; define an explicit circuit-breaker / request-shedding policy for extreme overload rather than letting queues grow unbounded.
Monitoring blind spot masks a slow SLA breach	Low	High	Dashboards and alerts defined in Section H, specifically on p95/p99 latency rather than only the average.

G. Deployment and Rollout Plan

Because the change is configuration-level only (Section E.6), it can be rolled out gradually rather than as a single cut-over, reducing the risk of an untested algorithm change causing an outage of its own.

Phase 1 – Shadow Mode

Run the Least Response Time routing logic in parallel with the existing Round Robin balancer without actually forwarding traffic through it, purely to log what decisions it would have made and validate that the EWMA values track real server performance sensibly.

Phase 2 – 5% Canary

Route a small, randomly-selected 5% slice of live traffic through the LRT balancer. Compare its p50/p95/p99 latency against the remaining 95% still on Round Robin, and roll back immediately if the canary's error rate or latency regresses instead of improving.

Phase 3 – 25% → 100% Ramp

If the canary is healthy for an agreed observation window, increase the LRT share to 25%, then 50%, then 100%, pausing at each step to confirm the metrics in Section H remain within target before proceeding.

Rollback Trigger Criteria

- p95 response time on the new balancer exceeds the old Round Robin baseline for more than a few consecutive monitoring windows.
- Any server's error rate rises noticeably above its pre-rollout baseline.
- The traffic-share fairness skew (Section E.5) exceeds an agreed operational limit, indicating a possible EWMA bug rather than intended behaviour.

H. Monitoring and Observability Plan

The engineering decision in Section E.6 is only as good as the monitoring that confirms it holds in production. The following metrics and alerts are proposed.

Metric	Why it matters	Suggested alert threshold
p50 / p95 / p99 response time (overall and per server)	The SLA in Section D is stated in terms of average and tail latency; per-server breakdowns catch a single struggling server before it affects everyone.	p95 overall > 2x the rolling 7-day baseline
Per-server traffic share vs. ideal (Section E.5)	Detects an EWMA bug or a mis-declared capacity weight before it becomes a latency incident.	Sustained deviation beyond an agreed fairness band
Per-server queue length / in-flight requests	Early warning of a server approaching saturation ($\rho \rightarrow 1$), before response time visibly degrades.	Queue length trending upward for 3+ consecutive intervals
Throughput (requests/second)	Confirms the new balancer is not silently dropping or rejecting requests.	Drop below the pre-rollout baseline
Health-check failure count per server	Feeds the reliability behaviour described in Section F.	2+ consecutive failures triggers removal from rotation

I. Glossary of Terms

Term	Definition
Load balancer	A component that sits in front of a pool of backend servers and decides, per request, which server should handle it.
SLA (Service Level Agreement)	An agreed performance target – here, an average response time of 2 seconds under normal load.
Latency / response time	The time from when a request arrives to when its response is returned, including any time spent waiting in a queue.
Tail latency (p95 / p99)	The response time below which 95% (or 99%) of requests complete; a better indicator of worst-case user experience than the average.
Throughput	The number of requests successfully completed per second.
Utilisation (ρ)	The fraction of a server's capacity currently in use, $\rho = \lambda/\mu$ (arrival rate over service rate).
Little's Law	$L = \lambda W$: the average number of requests in a system equals the arrival rate multiplied by the average time each spends in it.
EWMA (Exponentially-Weighted Moving Average)	A running average that weights recent observations more heavily than older ones, used here so the balancer 'forgets' stale performance data gradually.
Discrete-event simulation	A simulation technique that advances time only at the moments when something happens (an arrival, a completion), rather than in fixed small time steps.
Canary deployment	Releasing a change to a small fraction of traffic first, so problems are caught before they affect all users.
Thundering herd	A failure pattern where many requests pile onto the same resource at once, often after a reset event wipes out state that was previously spreading load evenly.

Appendix A: Simulation Source Code (Python)

The complete routing-and-simulation core is reproduced below for documentation and reproducibility, as required by Section E.10. This is the exact code used to produce every figure and table in Sections E.5 and Appendix B.

```
import simpy
import random
import statistics
import json

NUM_SERVERS = 4
WEIGHTS = [1, 1, 2, 3]
SIM_TIME = 2000
BASE_SERVICE_TIME = 1.2
SEEDS = [42, 7, 99, 123, 2024]
SCENARIOS = {
    "normal": 0.6, # mean inter-arrival seconds
    "peak": 0.3,   # flash-sale: ~2x arrival rate
}

def run_once(algorithm, mean_interarrival, seed):
    random.seed(seed)
    env = simpy.Environment()
    servers = [simpy.Resource(env, capacity=1) for _ in range(NUM_SERVERS)]
    queue_len = [0] * NUM_SERVERS
    ewma_rt = [BASE_SERVICE_TIME] * NUM_SERVERS # for least-response-time
    rr_counter = {"i": 0}
    wrr_state = {}
    response_times = []
    completed = [0]

    def service_time(idx):
        mean = BASE_SERVICE_TIME / WEIGHTS[idx]
        return random.expovariate(1 / mean)

    def pick_rr():
        i = rr_counter["i"] % NUM_SERVERS
        rr_counter["i"] += 1
        return i

    def pick_lc():
        return min(range(NUM_SERVERS), key=lambda i: queue_len[i])

    def pick_wrr():
        total = sum(WEIGHTS)
        for i in range(NUM_SERVERS):
            wrr_state[i] = wrr_state.get(i, 0) + WEIGHTS[i]
```



```

best = max(range(NUM_SERVERS), key=lambda i: wrt_state[i])
wrt_state[best] -= total
return best

def pick_lrt():
    # route to server with lowest exponentially-weighted average response time,
    # penalised by current queue length to avoid pile-up
    return min(range(NUM_SERVERS), key=lambda i: ewma_rt[i] * (1 + queue_len[i]))

picker = {
    "round_robin": pick_rr,
    "least_connections": pick_lc,
    "weighted_round_robin": pick_wrr,
    "least_response_time": pick_lrt,
}[algorithm]

def request_process(env, req_id):
    arrival = env.now
    idx = picker()
    queue_len[idx] += 1
    with servers[idx].request() as req:
        yield req
        st = service_time(idx)
        yield env.timeout(st)
    queue_len[idx] -= 1
    rt = env.now - arrival
    # update EWMA of this server's observed response time
    alpha = 0.2
    ewma_rt[idx] = alpha * rt + (1 - alpha) * ewma_rt[idx]
    response_times.append(rt)
    completed[0] += 1

def arrivals(env):
    while True:
        yield env.timeout(random.expovariate(1 / mean_interarrival))
        env.process(request_process(env, completed[0]))

env.process(arrivals(env))
env.run(until=SIM_TIME)

if not response_times:
    return None
avg_rt = statistics.mean(response_times)
p95_rt = statistics.quantiles(response_times, n=100)[94]
throughput = completed[0] / SIM_TIME
return {"avg_rt": avg_rt, "p95_rt": p95_rt, "throughput": throughput, "completed":

```

```


completed[0]}


ALGOS = ["round_robin", "least_connections", "weighted_round_robin",
        "least_response_time"]
full_results = {}
for scenario, mean_ia in SCENARIOS.items():
    full_results[scenario] = {}
    for algo in ALGOS:
        runs = [run_once(algo, mean_ia, s) for s in SEEDS]
        avg_rts = [r["avg_rt"] for r in runs]
        p95_rts = [r["p95_rt"] for r in runs]
        throughputs = [r["throughput"] for r in runs]
        completed = [r["completed"] for r in runs]
        full_results[scenario][algo] = {
            "avg_rt_mean": round(statistics.mean(avg_rts), 4),
            "avg_rt_std": round(statistics.stdev(avg_rts), 4),
            "p95_rt_mean": round(statistics.mean(p95_rts), 4),
            "p95_rt_std": round(statistics.stdev(p95_rts), 4),
            "throughput_mean": round(statistics.mean(throughputs), 4),
            "completed_mean": round(statistics.mean(completed), 1),
        }

print(json.dumps(full_results, indent=2))
with open("full_results.json", "w") as f:
    json.dump(full_results, f, indent=2)

```

IMPLEMENTATION:





Appendix B: Per-Seed Raw Simulation Results

The tables below give the individual result of every one of the 5 random seeds behind each mean ($\pm$ std) figure quoted in Section E.5, so the underlying run-to-run variability can be inspected directly rather than taken on trust.

Table B1: Normal load, per-seed results

Algorithm	Seed	Avg. Response Time	95th %ile	Throughput
Round Robin	42	1.190 s	4.265 s	1.655 req/s
Round Robin	7	1.169 s	4.243 s	1.701 req/s
Round Robin	99	1.087 s	3.798 s	1.674 req/s
Round Robin	123	1.143 s	4.131 s	1.659 req/s
Round Robin	2024	1.134 s	4.012 s	1.651 req/s
Least Connections	42	1.018 s	3.336 s	1.663 req/s
Least Connections	7	1.006 s	3.416 s	1.702 req/s
Least Connections	99	0.980 s	3.186 s	1.679 req/s
Least Connections	123	1.014 s	3.243 s	1.691 req/s
Least Connections	2024	1.030 s	3.287 s	1.637 req/s
Weighted Round Robin	42	0.811 s	2.600 s	1.680 req/s
Weighted Round Robin	7	0.771 s	2.574 s	1.685 req/s
Weighted Round Robin	99	0.767 s	2.505 s	1.676 req/s
Weighted Round Robin	123	0.788 s	2.593 s	1.681 req/s
Weighted Round Robin	2024	0.833 s	2.716 s	1.657 req/s
Least Response Time	42	0.668 s	2.045 s	1.639 req/s
Least Response Time	7	0.654 s	2.061 s	1.713 req/s
Least Response Time	99	0.650 s	1.949 s	1.675 req/s
Least Response Time	123	0.652 s	1.995 s	1.675 req/s
Least Response Time	2024	0.636 s	1.973 s	1.630 req/s

Table B2: Peak / flash-sale load, per-seed results

Algorithm	Seed	Avg. Response Time	95th %ile	Throughput
Round Robin	42	14.882 s	53.442 s	3.308 req/s
Round Robin	7	18.800 s	70.759 s	3.314 req/s
Round Robin	99	15.191 s	54.210 s	3.325 req/s
Round Robin	123	12.025 s	49.171 s	3.256 req/s
Round Robin	2024	13.727 s	54.585 s	3.216 req/s
Least Connections	42	1.100 s	3.854 s	3.292 req/s
Least Connections	7	1.096 s	3.696 s	3.350 req/s
Least Connections	99	1.143 s	3.835 s	3.381 req/s
Least Connections	123	1.117 s	3.711 s	3.337 req/s
Least Connections	2024	1.128 s	3.923 s	3.222 req/s

Weighted Round Robin	42	1.146 s	3.728 s	3.335 req/s
Weighted Round Robin	7	1.104 s	3.454 s	3.346 req/s
Weighted Round Robin	99	1.130 s	3.756 s	3.350 req/s
Weighted Round Robin	123	1.120 s	3.673 s	3.319 req/s
Weighted Round Robin	2024	1.143 s	3.771 s	3.224 req/s
Least Response Time	42	1.024 s	3.077 s	3.373 req/s
Least Response Time	7	0.944 s	2.871 s	3.346 req/s
Least Response Time	99	0.959 s	2.906 s	3.333 req/s
Least Response Time	123	0.959 s	2.890 s	3.341 req/s
Least Response Time	2024	0.934 s	2.749 s	3.203 req/s